

# converge-core

A FORMALLY-GROUNDED MULTI-AGENT RUNTIME

## The Problem with Agent Systems

---

Current multi-agent frameworks suffer from:

- **Non-determinism** — Same input, different outputs
- **No convergence guarantee** — Workflows can loop forever
- **Agent drift** — LLMs deviate from intended behavior
- **Untraceable decisions** — "Why did it do that?" has no answer
- **Partial failures** — Inconsistent state after errors

We can do better.

## The Converge Approach

---

A runtime built on **fixed-point semantics**:

```
Root Intent
  ↓
  Fan-out (agents propose)
  ↓
  Validate + enforce invariants
  ↓
  Serial commit (append-only Context)
  ↓
  Repeat until fixed point
  ↓
  Explainable result + audit trail
```

**Agents propose. The engine decides.**

# The 9 Axioms

MATHEMATICAL FOUNDATIONS

## Axioms 1-5

---

#	Axiom	Formula	Guarantee
1	<b>Monotonicity</b>	$\text{ctx} \subseteq \text{step}(\text{ctx})$	Facts never lost
2	<b>Determinism</b>	$\text{step}(\text{ctx}) = \text{step}(\text{ctx})$	Reproducible
3	<b>Idempotency</b>	$\text{agent}(\text{ctx}) = \text{agent}(\text{agent}(\text{ctx}))$	Safe retries
4	<b>Commutativity</b>	$a(b(\text{ctx})) = b(a(\text{ctx}))$	Order-independent
5	<b>Termination</b>	$\exists n: \text{step}^n(\text{ctx}) = \text{step}^{n+1}(\text{ctx})$	Always halts

## Axioms 6-9

---

#	Axiom	Formula	Guarantee
6	<b>Consistency</b>	$\neg\exists(f, \neg f) \in \text{ctx}$	No contradictions
7	<b>Starvation Freedom</b>	$\text{enabled}(a) \Rightarrow \diamond\text{runs}(a)$	Fair scheduling
8	<b>Confluence</b>	$\text{ctx}_1 \cup \text{ctx}_2 \rightarrow \text{ctx}^*$	Merges converge
9	<b>Observability</b>	$\forall \text{effect: } \text{logged(effect)}$	Full audit trail

These aren't guidelines — they're enforced by the type system.

# Core Architecture

TYPE-SAFE BY DESIGN

## The Context: Append-Only Truth

---

```
pub struct Context {
    facts: FactStore,          // Immutable, append-only
    clock: LamportClock,       // Logical ordering
    trace: ExecutionTrace,    // Full history
}

impl Context {
    pub fn derive(&self, new_facts: Vec<Fact>) -> Context {
        // Old context unchanged, new context returned
        Context {
            facts: self.facts.extend(new_facts),
            clock: self.clock.tick(),
            trace: self.trace.append(new_facts),
        }
    }
}
```

Immutability guarantees monotonicity.

## Facts: Typed, Versioned, Traceable

---

```
pub struct Fact {
    pub id: FactId,
    pub kind: FactKind,
    pub payload: Value,
    pub provenance: Provenance, // Who created this?
    pub timestamp: LamportTime,
    pub supersedes: Option<FactId>,
}

pub enum FactKind {
    Seed,           // Human-provided input
    Derived,        // Agent-computed
    ProposedFact,  // LLM suggestion (untrusted)
    Validated,      // Promoted from ProposedFact
}
```

## Agents: Pure Functions Over Context

---

```
pub trait Agent: Send + Sync {
    fn name(&self) -> &str;

    fn can_run(&self, ctx: &Context) -> bool;

    fn run(&self, ctx: &Context) -> AgentResult;
}

pub enum AgentResult {
    NoOp, // Nothing to do
    Propose(Vec<Fact>), // New facts to add
    Error(AgentError), // Recoverable failure
}
```

Agents are pure: `Context → AgentResult`

## The Engine: Fixed-Point Execution

---

```
impl Engine {
    pub fn run(&mut self, initial: Context) -> EngineResult {
        let mut ctx = initial;

        loop {
            let proposals = self.fan_out(&ctx);           // Parallel agent execution
            let validated = self.validate(proposals);      // Invariant checking
            let next = ctx.derive(validated);             // Immutable update

            if next == ctx {                            // Fixed point reached
                return EngineResult::Converged(next);
            }

            ctx = next;
            self.check_termination()?;
        }
    }
}
```

## LLM Integration: Trust Boundaries

---

```
pub struct LlmAgent {
    provider: Arc<dyn LlmProvider>,
    validator: Arc<dyn FactValidator>,
}

impl Agent for LlmAgent {
    fn run(&self, ctx: &Context) -> AgentResult {
        let response = self.provider.complete(ctx)?;

        // LLM output is NEVER directly trusted
        let proposals = response.facts.into_iter()
            .map(|f| f.as_proposed()) // Mark as ProposedFact
            .collect();

        AgentResult::Propose(proposals)
    }
}
```

**LLMs propose. Validators decide.**

## Validation: The Trust Layer

---

```
pub trait FactValidator: Send + Sync {
    fn validate(&self, fact: &Fact, ctx: &Context) -> ValidationResult;
}

pub enum ValidationResult {
    Accept,           // Promote to Validated
    Reject(RejectionReason), // Discard with reason
    NeedsReview,      // Human-in-the-loop
}

// Built-in validators
pub struct SchemaValidator;      // Type checking
pub struct InvariantValidator;   // Business rules
pub struct ConsistencyValidator; // No contradictions
```

# Key Features

WHAT MAKES CONVERGE-CORE DIFFERENT

## 1. Merkle-Based Integrity

---

```
pub struct FactStore {
    root: MerkleRoot,
    nodes: HashMap<FactId, MerkleNode>,
}

impl FactStore {
    pub fn verify(&self) -> bool {
        self.compute_root() == self.root
    }

    pub fn proof(&self, fact_id: &FactId) -> MerkleProof {
        // Generate proof that fact exists in store
    }
}
```

Tamper-evident audit trail. Cryptographic guarantees.

## 2. Parallel Agent Execution

---

```
impl Engine {
    fn fan_out(&self, ctx: &Context) -> Vec<AgentResult> {
        self.agents
            .par_iter() // Rayon parallel iterator
            .filter(|a| a.can_run(ctx))
            .map(|a| a.run(ctx))
            .collect()
    }
}
```

Commutativity axiom enables safe parallelism.

### 3. Time-Travel Debugging

---

```
impl ExecutionTrace {
    pub fn replay_to(&self, cycle: usize) -> Context {
        self.snapshots[..=cycle]
            .iter()
            .fold(Context::empty(), |ctx, facts| ctx.derive(facts))
    }

    pub fn diff(&self, from: usize, to: usize) -> Vec<Fact> {
        // Show exactly what changed between cycles
    }
}
```

Replay any execution. Debug any decision.

## 4. Pluggable LLM Providers

---

```
pub trait LlmProvider: Send + Sync {
    fn complete(&self, prompt: &Prompt) -> LlmResult<Response>;
    fn embed(&self, text: &str) -> LlmResult<Embedding>;
}

// Implementations
pub struct AnthropicProvider; // Claude
pub struct OpenAIProvider; // GPT-4
pub struct OllamaProvider; // Local models
pub struct MockProvider; // Testing
```

Swap providers without changing agent code.

## 5. Invariant DSL

---

```
invariant! {
    name: "no_negative_balance",
    description: "Account balance must never go negative",
    check: |ctx| {
        ctx.facts_of_type::<Balance>()
            .all(|b| b.amount >= 0)
    }
}

invariant! {
    name: "invoice_requires_delivery",
    description: "Cannot invoice without delivery proof",
    check: |ctx| {
        ctx.facts_of_type::<Invoice>()
            .all(|inv| ctx.has_delivery_for(&inv.work_id))
    }
}
```

# Performance

BUILT FOR PRODUCTION

## Benchmarks

---

Operation	Throughput	Latency (p99)
Fact insertion	100K/sec	0.1ms
Context derivation	50K/sec	0.2ms
Agent execution (pure)	10K/sec	1ms
Agent execution (LLM)	100/sec	500ms
Merkle verification	1M/sec	0.01ms

Measured on M2 MacBook Pro, single-threaded

## Memory Model

---

```
Context Size vs Facts:  
1K facts → ~100 KB  
10K facts → ~1 MB  
100K facts → ~10 MB  
1M facts → ~100 MB
```

Derivation is  $O(n)$  where  $n$  = new facts only  
Full context never copied, only extended

Structural sharing via immutable data structures.

# Getting Started

FROM ZERO TO RUNNING

## Installation

---

```
# Cargo.toml
[dependencies]
converge-core = "0.1"
converge-provider = "0.1"  # LLM providers
tokio = { version = "1", features = ["full"] }
```

```
# Or use the CLI
cargo install converge-cli
converge new my-project
```

## Hello, Converge

---

```
use converge_core::{Engine, Context, Agent, Fact};

struct HelloAgent;

impl Agent for HelloAgent {
    fn name(&self) -> &str { "hello" }
    fn can_run(&self, ctx: &Context) -> bool {
        !ctx.has_fact_of_type::<Greeting>()
    }
    fn run(&self, _ctx: &Context) -> AgentResult {
        AgentResult::Propose(vec![
            Fact::derived("greeting", Greeting { message: "Hello, Converge!" })
        ])
    }
}

fn main() {
    let mut engine = Engine::new();
    engine.register(HelloAgent);
    let result = engine.run(Context::empty());
    println!("{:?}", result);
}
```

## Project Structure

---

```
my-project/
├── Cargo.toml
└── src/
    ├── main.rs
    ├── agents/          # Your agent implementations
    │   └── mod.rs
    │   └── invoice.rs
    ├── facts/           # Domain fact types
    │   └── mod.rs
    │   └── money.rs
    └── invariants/      # Business rules
        └── mod.rs
    └── specs/           # Converge Truths (.truth files)
        └── money.truth
```

## Documentation & Resources

---

- **Docs:** [docs.rs/converge-core](https://docs.rs/converge-core)
- **Examples:** [github.com/kpernyer/converge-core/examples](https://github.com/kpernyer/converge-core/examples)
- **Discord:** [discord.gg/converge](https://discord.gg/converge)
- **Blog:** [converge.zone/signals](https://converge.zone/signals)

## Let's Talk

---



**Kenneth Pernyer** · Founder

[LinkedIn](#)

[Twitter](#)

[GitHub](#)